

YUPENG SU

UC Santa Barbara | yupengsu@ucsb.edu | yupengsu.github.io | linkedin.com/in/yupeng-su

Research Interests

Efficient Pre-training/Post-training Optimization for Large Language Models: focusing on quantization, sparsification, and low-rank decomposition techniques to improve training and inference efficiency

High-Performance Computing Infrastructure on GPU/FPGA Architectures: including compiler optimization and GPU/FPGA kernel adaptation for high-performance, energy-efficient inference and practical edge deployment

Education

- Adviser: Prof. Zheng Zhang (ECE, UGSD)

Southern University of Science and Technology, B.E. in Microelectronics – China
Adviser: Prof. Han Yu (SUSTech)

- Advisor: Prof. Hao Yu (SME, SUSTech).
 - CGA/GPA: 3.93/4.0, with 178 units completed.
 - Awards: Nominee of Top Ten Graduates of SUSTech, Top Ten Graduates of the College of Engineering, SUSTech
 - Dissertation: *Enhanced Mix-Precision Post-Training Quantization for Large Language Models: Block Awareness and Outlier Identification.*

Publications

APTQ: Attention-aware post-training mixed-precision quantization for large language models

Ziyi Guan, Hantao Huang, **Yupeng Su**, Hong Huang, Ngai Wong, Hao Yu

[10.1145/3649329.3658498](https://doi.org/10.1145/3649329.3658498) (IEEE/ACM DAC, 2024)

LLM-Barber: Block-Aware Rebuilder for Sparsity Mask in One-Shot for Large Language Models

Yupeng Su, Ziyi Guan, Xiaoqun Liu, Tianlai Jin, Dongkuan Wu, Chenfei Chen, Graziano Chesi, Ngai Wong, Hao Yu
10.1109/ICCAD66269.2025.11240704 (IEEE/ACM ICCAD, 2025)

EdgeLLM: A Highly Efficient CPU-FPGA Heterogeneous Edge Accelerator for Large Language Models

Mingqiang Huang, Ao Shen, Kai Li, Haoxiang Peng, Boyu Li, Yupeng Su, Hao Yu

10.1109/TCSI.2025.3546256 (IEEE TCAS I: Regular Papers)

Quantization Meets Reasoning: Exploring Low-Bit Quantization Degradation for Mathematical Reasoning

Zhen Li*, Yupeng Su*, Runming Yang, Congkai Xie, Zheng Wang, Zhongwei Xie, Ngai Wong, Hongxia Yang

arxiv.org/abs/2501.03035 (Under Review, *Equal Contribution)

PTOTP: Post-Training Quantization to Trit-Planes for Large Language Models

He Xiao, Runming Yang, Qingyao Yang, Wendong Xu, Zhen Li, **Yupeng Su**, Zhengwu Liu, Hongxia Yang, Ngai Wong
arxiv.org/abs/2509.16989 (Under Review)

Experience

Student Research Assistant, Prof. Hongxia Yang's Lab @ PolyU – Hongkong, China Nov 2024 – May 2025
• Investigated the effects of compression on reasoning and explored data-driven strategies for restoration.

- Investigated the effects of compression on reasoning and explored data-driven strategies for restoration.

Student Research Assistant, Next-Gen AI (NgaI) Lab @ HKU - Hong Kong, China Aug 2024 - Feb 2025

- Contributed to the lightweight networks for efficient inference under specified computing architectures.

Summer Intern, High Performance Integrated Circuit Design Lab @ SUSTech – China June 2024 – Sept 2024

- Published EdgeLM about the project and results

- Published [EdgeLEM](#) about the project and results